

✓ VGG16 Architecture Implementation

Invidia-smi

Sat Aug 24 05:36:54 2024

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI 535.104.05				Driver Version: 535.104.05			CUDA Version: 12.2		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU Name		Persistence-M		Bus-Id		Disp.A		Volatile Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util		Compute M.	MIG M.
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
0	Tesla T4		Off	00000000:00:04:0		Off			0
N/A	48C	P0	27W / 70W	14091MiB / 15360MiB		0%		Default	N/A
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
Processes:									
GPU	GI	CI	PID	Type	Process name			GPU Memory	
	ID	ID						Usage	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os

import cv2
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

from keras.models import Sequential
from keras.layers import Flatten, Dense, Input, Conv2D, MaxPool2D, Dropout
from keras.preprocessing import image
from keras.applications.imagenet_utils import preprocess_input
from keras.optimizers import SGD
```

Suggested code may be subject to a license | AAAAaron/Object-Detection-Learn |
def VGG16(input_tensor=None, classes=2):

```
    model = Sequential()

    # Block1
    model.add(Conv2D(64, (3,3), activation='relu', padding='same', name='block1_conv1', input_shape=(300,300,3)))
    model.add(Conv2D(64, (3,3), activation='relu', padding='same', name='block1_conv2'))
    model.add(MaxPool2D((2,2), strides=(2,2), name='block1_maxpool'))

    # Block2
    model.add(Conv2D(128, (3,3), activation='relu', padding='same', name='block2_conv1'))
    model.add(Conv2D(128, (3,3), activation='relu', padding='same', name='block2_conv2'))
    model.add(MaxPool2D((2,2), strides=(2,2), name='block2_maxpool'))

    # Block3
    model.add(Conv2D(256, (3,3), activation='relu', padding='same', name='block3_conv1'))
    model.add(Conv2D(256, (3,3), activation='relu', padding='same', name='block3_conv2'))
    model.add(Conv2D(256, (3,3), activation='relu', padding='same', name='block3_conv3'))
    model.add(MaxPool2D((2,2), strides=(2,2), name='block3_maxpool'))

    # Block4
    model.add(Conv2D(512, (3,3), activation='relu', padding='same', name='block4_conv1'))
    model.add(Conv2D(512, (3,3), activation='relu', padding='same', name='block4_conv2'))
    model.add(Conv2D(512, (3,3), activation='relu', padding='same', name='block4_conv3'))
    model.add(MaxPool2D((2,2), strides=(2,2), name='block4_maxpool'))

    # Block5
    model.add(Conv2D(512, (3,3), activation='relu', padding='same', name='block5_conv1'))
    model.add(Conv2D(512, (3,3), activation='relu', padding='same', name='block5_conv2'))
    model.add(Conv2D(512, (3,3), activation='relu', padding='same', name='block5_conv3'))
    model.add(MaxPool2D((2,2), strides=(2,2), name='block5_maxpool'))
```

```
# Block6
model.add(Flatten(name="flatten"))
model.add(Dense(4096, activation='relu', name='fc1'))
model.add(Dropout(0.5))
model.add(Dense(4096, activation='relu', name='fc2'))
model.add(Dropout(0.5))
model.add(Dense(classes, activation='softmax', name='prediction'))

return model
```

```
model = VGG16(classes=2)
```

```
model.summary()
```

→ Model: "sequential_3"

Layer (type)	Output Shape	Param #
block1_conv1 (Conv2D)	(None, 300, 300, 64)	1,792
block1_conv2 (Conv2D)	(None, 300, 300, 64)	36,928
block1_maxpool (MaxPooling2D)	(None, 150, 150, 64)	0
block2_conv1 (Conv2D)	(None, 150, 150, 128)	73,856
block2_conv2 (Conv2D)	(None, 150, 150, 128)	147,584
block2_maxpool (MaxPooling2D)	(None, 75, 75, 128)	0
block3_conv1 (Conv2D)	(None, 75, 75, 256)	295,168
block3_conv2 (Conv2D)	(None, 75, 75, 256)	590,080
block3_conv3 (Conv2D)	(None, 75, 75, 256)	590,080
block3_maxpool (MaxPooling2D)	(None, 37, 37, 256)	0
block4_conv1 (Conv2D)	(None, 37, 37, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 37, 37, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 37, 37, 512)	2,359,808
block4_maxpool (MaxPooling2D)	(None, 18, 18, 512)	0
block5_conv1 (Conv2D)	(None, 18, 18, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 18, 18, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 18, 18, 512)	2,359,808
block5_maxpool (MaxPooling2D)	(None, 9, 9, 512)	0
flatten (Flatten)	(None, 41472)	0
fc1 (Dense)	(None, 4096)	169,873,408
dropout_8 (Dropout)	(None, 4096)	0
fc2 (Dense)	(None, 4096)	16,781,312
dropout_9 (Dropout)	(None, 4096)	0
prediction (Dense)	(None, 2)	8,194

Total params: 201,377,602 (768.19 MB)
Trainable params: 201,377,602 (768.19 MB)

```
opt = SGD(learning_rate=0.01)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics = ['accuracy'])
```

```
training_ds_path = '/content/drive/MyDrive/datasets/img-clsf/dataset/train'
testing_ds_path = '/content/drive/MyDrive/datasets/img-clsf/dataset/test'
```

```
cats_n_dogs = os.listdir(training_ds_path)
print(cats_n_dogs)
```

→ ['dogs', 'cats']

```
arr_cats_dogs = []

for item in cats_n_dogs:
    #print(item)
    all_c_d = os.listdir(training_ds_path + "/" + item)
    #print(training_ds_path + "/" + item)

    for cat_dog in all_c_d:
        arr_cats_dogs.append((item, str(training_ds_path+ "/" + item) + '/' + cat_dog))

print(arr_cats_dogs)
```

→ [['dogs', '/content/drive/MyDrive/datasets/img-clsf/dataset/train/dogs/dog_293.jpg'), ('dogs', '/content/drive/MyDrive/datasets/img-clsf

```
cats_n_dogs_df = pd.DataFrame(data=arr_cats_dogs, columns=['cats_dogs','image'])
print(cats_n_dogs_df)
```

→

	cats_dogs	image
0	dogs	/content/drive/MyDrive/datasets/img-clsf/datas...
1	dogs	/content/drive/MyDrive/datasets/img-clsf/datas...
2	dogs	/content/drive/MyDrive/datasets/img-clsf/datas...
3	dogs	/content/drive/MyDrive/datasets/img-clsf/datas...
4	dogs	/content/drive/MyDrive/datasets/img-clsf/datas...
..
552	cats	/content/drive/MyDrive/datasets/img-clsf/datas...
553	cats	/content/drive/MyDrive/datasets/img-clsf/datas...
554	cats	/content/drive/MyDrive/datasets/img-clsf/datas...
555	cats	/content/drive/MyDrive/datasets/img-clsf/datas...
556	cats	/content/drive/MyDrive/datasets/img-clsf/datas...

[557 rows x 2 columns]

```
print("Total number of images: ", len(arr_cats_dogs))
cats_n_dogs_df['cats_dogs'].value_counts()
```

→ Total number of images: 557

	count
cats_dogs	
cats	279
dogs	278

dtype: int64

```
#path = './dataset/train/'

im_size = 300

images = []
labels = []

for indx in cats_n_dogs:
    data_path = training_ds_path + "/" + str(indx)

    image_files = [image for image in os.listdir(data_path)]

    for img_file in image_files:
        img = cv2.imread(data_path + '/' + img_file)
        img = cv2.resize(img, (im_size, im_size))
        images.append(img)
        labels.append(indx)
```

```
images = np.array(images)
images = images.astype('float32') / 255.0
images.shape
```

→ (557, 300, 300, 3)

```
le = LabelEncoder()
y = le.fit_transform(y)
print(y)
le.classes_
```

```
y = y.reshape(-1,1)

ohe = OneHotEncoder(sparse_output=False)
y = ohe.fit_transform(y)
y.shape
```

```
y
array([[0., 1.],
       [0., 1.],
       [0., 1.],
       ...,
       [1., 0.],
       [1., 0.],
       [1., 0.]])
```

```

Epoch 1/15
18/18 ————— 22s 1s/step - accuracy: 0.5123 - loss: 0.6932
Epoch 2/15
18/18 ————— 35s 832ms/step - accuracy: 0.5080 - loss: 0.6931
Epoch 3/15
18/18 ————— 21s 841ms/step - accuracy: 0.4550 - loss: 0.6933
Epoch 4/15
18/18 ————— 21s 847ms/step - accuracy: 0.4962 - loss: 0.6933
Epoch 5/15
18/18 ————— 20s 850ms/step - accuracy: 0.5297 - loss: 0.6929
Epoch 6/15
18/18 ————— 16s 858ms/step - accuracy: 0.5335 - loss: 0.6927
Epoch 7/15
18/18 ————— 21s 862ms/step - accuracy: 0.5351 - loss: 0.6929
Epoch 8/15
18/18 ————— 20s 864ms/step - accuracy: 0.5359 - loss: 0.6929
Epoch 9/15
18/18 ————— 21s 869ms/step - accuracy: 0.5088 - loss: 0.6930
Epoch 10/15
18/18 ————— 21s 871ms/step - accuracy: 0.4953 - loss: 0.6932
Epoch 11/15
18/18 ————— 16s 880ms/step - accuracy: 0.5200 - loss: 0.6930
Epoch 12/15
18/18 ————— 21s 882ms/step - accuracy: 0.5019 - loss: 0.6931
Epoch 13/15
18/18 ————— 20s 881ms/step - accuracy: 0.5360 - loss: 0.6930
Epoch 14/15
18/18 ————— 16s 885ms/step - accuracy: 0.5158 - loss: 0.6929
Epoch 15/15
18/18 ————— 20s 885ms/step - accuracy: 0.5149 - loss: 0.6930
<keras.src.callbacks.history.History at 0x7e35e1a31000>

```

```
from matplotlib.pyplot import imread
from matplotlib.pyplot import imshow
```

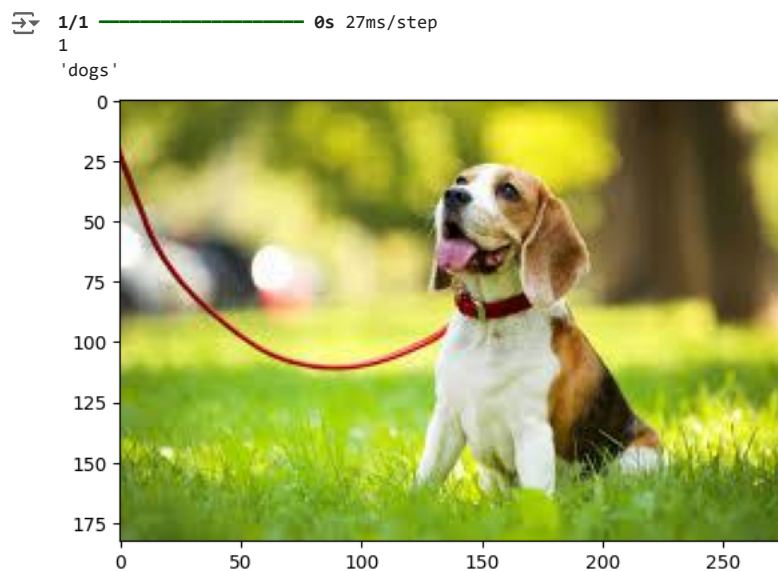
```
def PredictImage(image_path):
    img = image.load_img(image_path, target_size=(300, 300))

    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)
    x = preprocess_input(x)
    my_image = imread(image_path)
    imshow(my_image)

    values = model.predict(x)
    print( np.argmax(values))

    return le.classes_[np.argmax(values)]
```

```
PredictImage('/content/drive/MyDrive/datasets/img-clsf/dataset/test/dogs/dog_354.jpg')
```



```
PredictImage('/content/drive/MyDrive/datasets/img-clsf/dataset/test/cats/cat_162.jpg')
```

